

1. Good Evening

2. Lecture begins at 9:05

3. Topic - Indexes &
Backlog

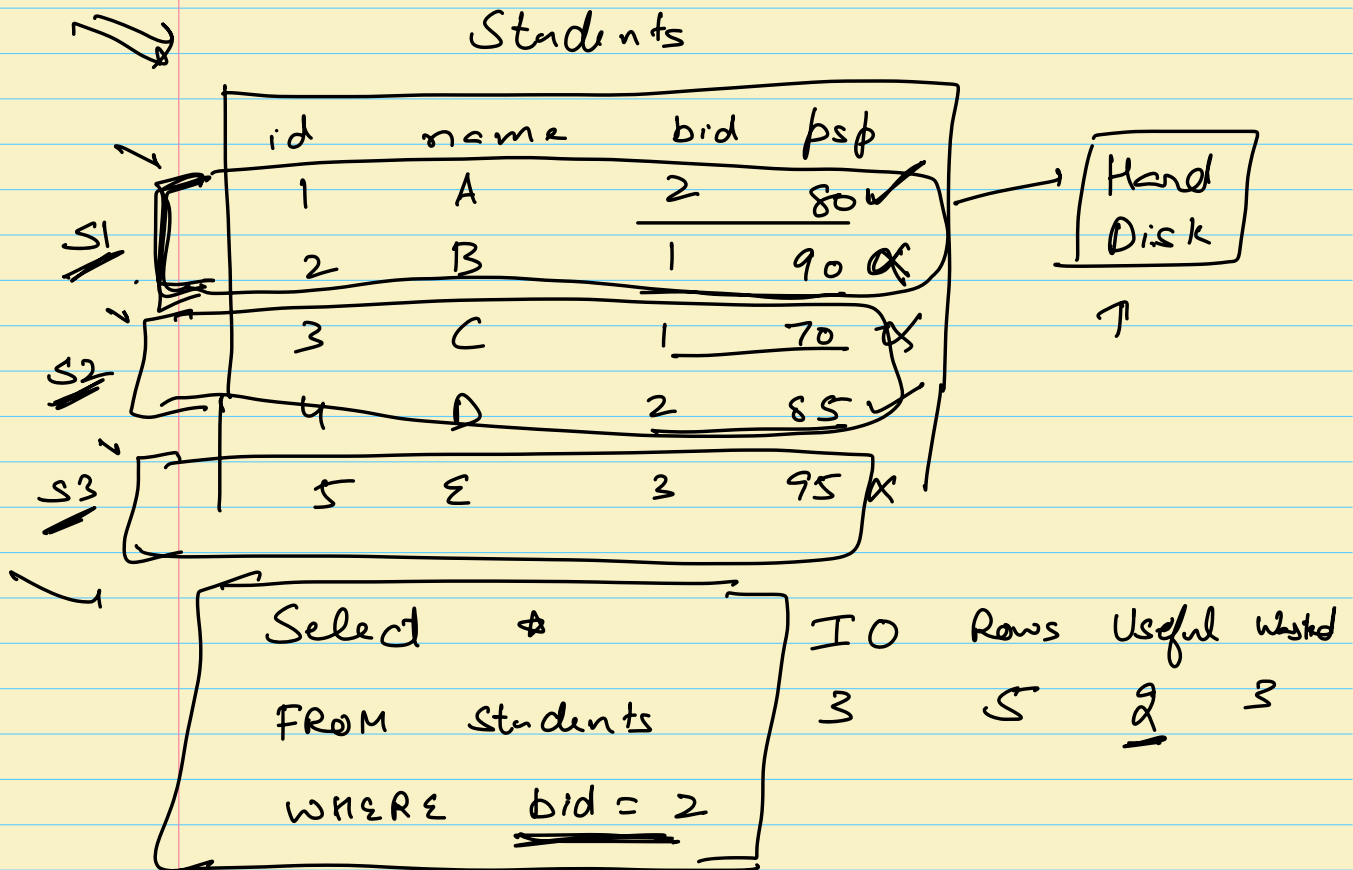
Agenda

1. Indexes

2. Backlog → [Data Types.
Built-in Fns.
Serializable → ✓
Deadlocks]

Intro of Indexes.

Consider the table.



CPU — 3 GHz $[3 \times 10^9 \text{ ops/second}]$

[Registers >> Cache >> RAM >> Hard Disk]

↑ ↑ ↔ ↑

1. IO operation will bring data from hard-disk to RAM, segment by segment.
2. Once the data is in RAM, it is analysed by CPU & rows are selected/rejected to prepare result.

Approach 1

Students					
id	name	bid	ps p		
1	A	2	80	}	s ¹
2	B	1	90		
3	C	1	85	}	s ²
⋮		3			
⋮					

↓
1000 Rows

5 rows don't
bid = 2

1000 ... 2 —


IO	Rows	Useful	Wasted
<u>500</u>	<u>1000</u>	<u>5</u>	<u>995</u>

Query :

Select *
FROM students
WHERE id = 5 ✓

Fact → Data of a table is sorted
by PK on the hard-disk

	id	name	bid	psp	
1000 Rows	1	A	2	—	S ₁ ✓
	2	B	1	—	
	3	C	1	—	S ₂ ✓
	4	D	2	—	
	5	E	5	—	S ₃ ✓
	⋮				




IO	Rows	Useful	visited.
3	6	1	5

SUMMARY → When clause on Pk will always work faster so data is sorted by Pk & Pk is unique also. We will stop reading future segments when we will find desired row.

Approach 3 : Lookup Table / Index

→ For every table, MySQL creates a lookup table or index [stored in RAM as well as hard-disk] for the Pk.



6. F 2

6. 53

DSA = BTree

Select *
FROM Students
WHERE id = 5

1. Check lookup for segment.
2. Read segment
3. Process.

IO
1.

Rows
2

Useful
1

Wasted
1

Question : How fast will the search
be in lookup table?

DSA used to implement Index = BTree

Case 1	Case 2	Case 3
id	id — sorted	id — sorted Index

Select *

FROM Students

WHERE bid = 2

Students

id	name	bid	psp	
1	A	2	—	S1
2	B	1	—	
3	C	1	—	S2
4	D	3	—	
5	E	2	—	S3
6	F	1	—	
7	G	3	—	S4

* Pk has index by default.

* We can create index for other columns as required.

↓

	key(bid)	val (List of symbols)
1	1	S1, S2, S3
1	<u>2</u>	S1, S3
	3	S2, S4

BTree

IO	Rows	Useful	Wasted
2	4	2	2

- Pros & Cons
- Guidelines
- Composite Indexes

BREAK = 10:18 to 10:25

1. Demo
- show
 - CREATE
 - DROP
 - explain

✓


2. DSA for index ✓

3. Pros & Cons

4. Guidelines.

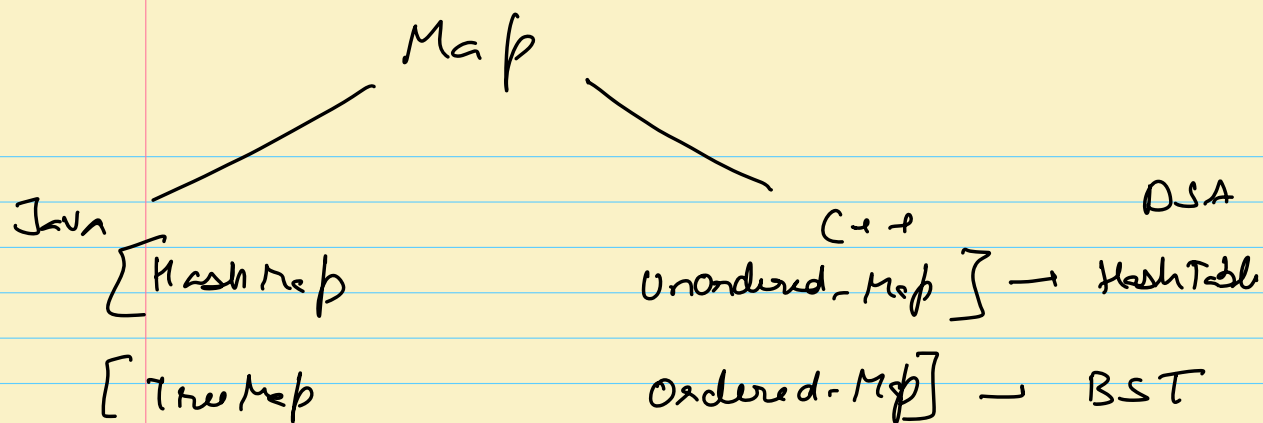
5. Indexes on String

6. Composite Indexes → (A1) (A3)





DSA FOR Index

key	value
1	S1
2	S1
—	—
—	—



HashMap → Avg constant

Balanced BST → $O(\log n)$



$\log_2 n$

Faster?

$$\log_2 n$$

\ll

$$\log_{10} n$$


$$\log_2 1024$$

10 operations.

$$\log_{10} 1024$$

3 comparisons.

B Tree



$$T_{\text{tree}} = \log_{26} n$$

Algorithm = \log_n
degree

B Trees { degree can be as high }
as 2000 }

\log_n
2000

8×10^9 rows

\log_n
2000

id
1
2
.
.
.
 8×10^9 rows

$\log_{2000} 8 \times 10^9$

= 3 comparisons

1. Indexes use B-Trees.
2. B-Trees are similar to BSTs
3. The Performance of B-Tree is $\log n$ _{degree}

⇒ How to implement using B-Trees for DBMS

★ PROS & CONS of Indexes.

PROS

1. Fast search ^{oo}_o of optimised reads.
We read only relevant segments, we decrease the number of non-useful rows read.

CONS

1. Extra space per lookup.
2. INSERT, Delete, Update might become slower.

Students

id	name	bid	psp
1	A	2	80
2	B	1	90
3	C	1	75
4	D	2	80
5	E	3	90
6	F	1	75
7	G	2	80

S1

S2

S3



Join-id

id	segment
1	S1
2	S1
3	S2
4	S2
5	S3
6	S3

Join-Name

n	segment
A	S1
B	S1
C	S2
D	S2
E	S3
F	S3

Join-bid

bid	segment
2	S1, S2
1	S1, S2, S3
3	S3


Join-psp

psp	segment
-----	---------

GUIDELINES

1. Don't create index while creating tables [let the pk index get created]
2. Analyse the traffic on db.

SQL Profiler.



3. Create index for columns, which are frequently used by where clause [as per data logged by SQL Profiler].

Indexes on String

Students

id	name	bid	psp
1	<u>Sumeet</u>	—	—
2	<u>Maganh</u>	—	—
3	<u>Naman</u>	—	—
4	<u>Megha</u>	—	—
5	<u>Smriti</u>	—	—
6	<u>Nayan</u>	—	—
7	<u>Mukesh</u>	—	—
8	<u>Sudheer</u>	—	—

S1. (rows 1, 2)
 S2. (rows 3, 4)
 S3. (rows 5, 6)
 S4. (rows 7, 8)

Sudheer

Case 1 : No index.

IO	Rows	Useful	Wasted.
4	8	1	7

Case 2 : Index on name column

IO	Rows	Useful	Wasted.
1	2	1	1

Sumeet	S1
Maganh	S1
Naman.	S2
Megha	S2
Smriti	S3
Nayan	S3
Mukesh	S4

| Sudheer S4 |

Case 3: Index on 1st char of name.

IO	Rows	Useful	Wasted.	
3	6	1.	5	<div style="border: 1px solid black; padding: 5px; display: inline-block;"> <u>S</u> : <u>S1, S3, S4</u> <u>M</u> : S1, S2, S4 <u>N</u> : S2, S3 </div>

Case 3: Index of 1st 2 char

IO	Rows	Useful	Wasted	
2	4	1	3	<div style="border: 1px solid black; padding: 5px; display: inline-block;"> Su → <u>S1, S4</u> Ma → S1 Na → S2, S3 Me → S2 Sm → S3 Mu → S4 </div>

Case 4: Index on 1st 3 chars

IO

Row

Useful wasted

1

2

1

1

Sum

— S1

Mag

— S1

Num

— S2

Mag

— S2

Sum

— S3

Mag



— S3

Mark


— S4

Sum

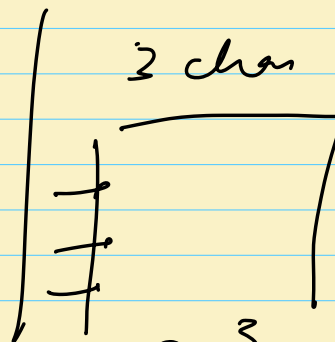
— S4



2^6

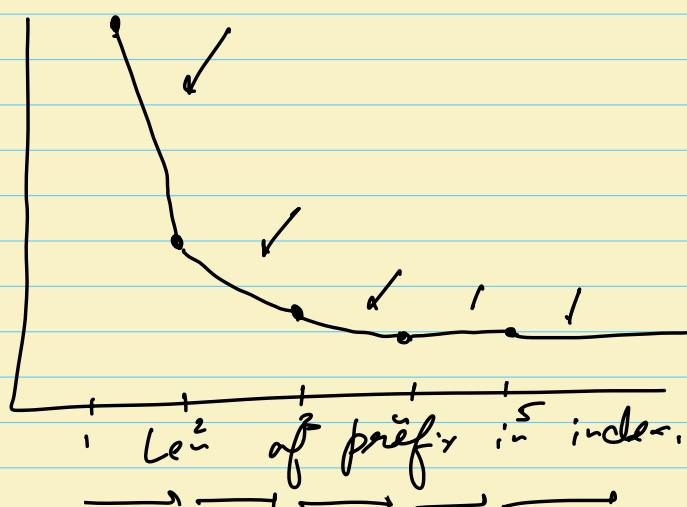


2^6^2



2^6^3

IO operations




Backlog

+ PROCEDURES

+ UDFs.

1. Data Types
2. Builtin Fns.
3. SERIALIZABLE + Demo + Deadlocks
4. COMPOSITE Indexes
5. Full-text SEARCH

SUB-QUERY + VIEWS



Robert Vieira

- Beginning MS SQL Server 2005
- MS SQL Server 2012 for Advanced Programmers

